


SMT-RAT

Version 0.4.0

Satisfiability-Modulo-Theories Real Arithmetic Toolbox

Manual

Florian Corzilius, Ulrich Loup, Sebastian Junges, and Erika Ábrahám



Contents

1	Introduction	3
2	Installation	5
2.1	Requirements	5
2.2	Building, installing and uninstalling SMT-RAT	5
3	General framework of SMT-RAT	7
3.1	The Formula class	7
3.2	The Module class	7
3.3	The Manager and the Strategy class	8
4	Embedding of a theory solver into an SMT solver	9
4.1	Interfaces of a theory solver composed with SMT-RAT	9
4.2	Example: Embedding a SMT-RAT composition into OpenSMT	10
5	Implementing further modules	17
5.1	Quickstart	17
5.2	Members of a modul	18
5.3	Virtual interfaces	19
5.3.1	Informing about a constraint	20
5.3.2	Adding a constraint	20
5.3.3	Checking for consistency	20
5.3.4	Pushing a backtrack point	21
5.3.5	Popping a backtrack point	21
5.4	Running the backend module	22
6	Composing a solver	25
6.1	The strategy	25
6.2	Specifying a strategy with the GUI	26

1 Introduction

This manual describes **SMT-RAT**, a toolbox which is specifically tailored to be used by an SMT solver in order to solve non-linear real arithmetic (NRA) efficiently. NRA is an important but hard-to-solve theory and only a fragment of it can be handled by some of the currently existing SMT solvers. The toolbox consists of four modules, implementing the virtual substitution method, the cylindrical algebraic decomposition method, a groebner bases simplifier and a general simplifier, respectively. The intention of the toolbox is that an SMT-compliant theory solver can be achieved by composing these modules according to a user defined strategy and with the goal to exploit their advantages. Furthermore, it supports the addition of further modules implementing other methods dealing with real arithmetic.

2 Installation

2.1 Requirements

SMT-RAT has been successfully compiled and tested under Linux. For its configuration we use `cmake`, which can be found on <http://www.cmake.org/>, and for its compilation we use `g++` (version 4.7 or higher), available on <http://gcc.gnu.org/>. The configuration settings can be changed with the command line interface `ccmake`. The arithmetic calculations within SMT-RAT use the C++ library `GiNaC`, which can be found on <http://www.ginac.de/>. The algebraic operations and data structures within SMT-RAT are based on the C++ library `GiNaCRA` available at <http://ginacra.sourceforge.net/>. Optional, there is a graphical user interface for the composition of a solver, which needs `java` (version 1.7 or higher). Note, that we use C++11. Summarizing, you need the following packages before installation:

- `g++` (version 4.7 or higher)
- `cmake`
- `ccmake`
- `GiNaC`
- `GiNaCRA`
- `java` (version 1.7 or higher)

2.2 Building, installing and uninstalling SMT-RAT

You can download SMT-RAT from <http://smtrat.sourceforge.net/> and install it the following way:

1. Unpack the package:

```
tar xvzf smtrat-*.tar.gz
```

2. Create a directory (e.g. `build`) that will contain the object files and the executables, and change into it:

```
mkdir build && cd build
```

3. Configure:

2 *Installation*

```
cmake ..
```

4. Build:

```
make
```

5. Install:

```
make install
```

6. Uninstall:

```
xargs rm < install_manifest.txt
```

```
make clean
```

More information can be found in the README file, which can be found in **SMT-RAT** directory.

3 General framework of SMT-RAT

Our toolbox has a modular C++ class design which can be used to compose NRA theory solvers for an SMT embedding in a *dynamic* and *hierarchical* fashion. Our NRA theory solvers are instances of **Manager**, which offers an interface to communicate with the environment and which coordinates the satisfiability check according to a user-defined **Strategy**. Such a strategy combines basic NRA theory solver modules, derived from **Module**. Figure 3.1 shows an example configuration. Moreover, a **Java**-based graphical user interface (GUI) can be used for an intuitive and user-friendly specification of strategies (and the automatic generation of a corresponding **Strategy** class). Next, we briefly describe these concepts.

SMT-RAT is a C++ library consisting of a collection of SMT-compliant modules which can be combined to (1) a theory solver in order to extend the supported logics of an existing SMT solver by non-linear real arithmetic and (2) to SMT solver in which these modules can be (and most of them are) integrated to tackle RCF. The latter is intended to be a testing environment for the development of SMT-compliant theory solvers, as the one presented in this paper. **SMT-RAT** defines three types of components (see Appendix ??): *manager*, *strategy* and *module*. In the following we first describe the functionality of a module and show how the manager composes different modules according to a strategy to a solver.

3.1 The Formula class

Formula instances contain, besides a sequence of NRA constraints, a bitvector storing some information about the problem and the history of its check. E.g., there is a bit which is 1 if some of the constraints are equations. Also for each module there is a bit which is 1 if the module was already invoked on the given problem. Such information can be used to specify conditions under which a procedure should be invoked for a certain problem.

3.2 The Module class

A module m contains a set of formulas, called its *set of received formulas*, we denote by $C_{rcv}(m)$. The main procedure of a module is `isConsistent()` and either decides whether C_{rcv} is satisfiable or not returning `sat` or `unsat`, respectively, or returns `unknown`. Note, that a set of formulas is semantically defined by their conjunction. We can manipulate the set of received formulas by adding (removing) formulas φ to (from) it with `add(φ)` (`remove(φ)`). Since in the SMT embedding C_{rcv} is usually changed

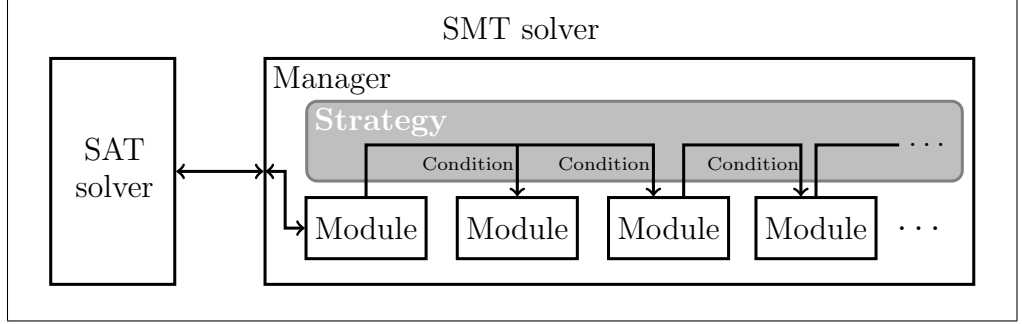


Figure 3.1: A snapshot of an SMT-RAT composition embedded in an SMT solver.

between two consecutive `isConsistent()` calls only by adding/removing constraints, the solver's performance can be significantly improved if the modules can make use of the results of previous checks, that is they support *incrementality* and *backtracking*. In case that the module determines the unsatisfiability of C_{rcv} , it is expected to compute at least one preferably small *infeasible subset* $C_{inf} \subseteq C_{rcv}$. Moreover, a module has the possibility to name lemmas, which are formulas being tautologies, that is they hold for all assignments of the variables occurring in them. These lemmas should encapsulate information which can be extracted from a module's internal state and propagated among other SMT-RAT modules. Furthermore, SMT-RAT provides the feature that a module itself can ask other modules for the satisfiability of the set C_{pas} of formulas, called its *set of passed formulas*, using the procedure `runBackends()` which is controlled by the manager.

3.3 The Manager and the Strategy class

A *strategy* is a directed tree $T := (V, E)$ with a set V of module instances as nodes and $E \subseteq V \times \Omega \times V$, where Ω is a set of conditions. The *manager* contains the strategy and the input formula C_{input} , either received by a prefixed solver or parsed from an example file. Furthermore, it maintains the allocation of module instances as follows. Initially, the manager calls the method `isConsistent()` of the module instance given by the root of the strategy with $C_{rcv} = C_{input}$ being the set of received formulas of this module. Whenever a module instance $m \in V$ calls `runBackends()`, with C_{pas} being its set of passed formulas, the manager calls `isConsistent()` of each module instance $m' \in V$ with $C_{rcv}' = C_{pas}$ being its set of received formulas, for which an edge $(m, \omega, m') \in E$ exists such that ω holds for C_{pas} , and passes the results back to m . Furthermore, it also passes back the infeasible subsets and lemmas provided by the invoked modules. The module m can now benefit in its solving and reasoning process from this shared information.

4 Embedding of a theory solver into an SMT solver

In the following we assume, that either a composed theory solver according to Chapter 6 or the composition we provide, i.e. the `NRATSSolver`, is used. They only differ in the strategy they implement and provide the same interfaces. We give a short description of these interfaces and give an example by embedding it into `OpenSMT`.

4.1 Interfaces of a theory solver composed with SMT-RAT

An SMT solver can use the following interfaces:

- `bool inform(const string& _constraint, bool _infix)`

Informs the theory solver about the existence of the given constraint in form of a string. The second argument is a flag, which indicates whether the given constraint is written in infix or prefix notation.

- `bool addConstraint`
(
 `const string& _constraint,`
 `const bool _infix,`
 `const bool _polarity`
)

Adds the given constraint in form of a string to the theory solver. The second argument is a flag, which indicates whether the given constraint is written in infix or prefix notation. The last argument is again a flag which is `true` if the constraint has to hold in the given form and `false` if its inversion has to hold. The inversion operator for a constraint simply changes its relation symbol in the following way:

$$\begin{array}{lcl} = & \mapsto & \neq \\ \neq & \mapsto & = \\ \leq & \mapsto & > \\ \geq & \mapsto & < \\ < & \mapsto & \geq \\ > & \mapsto & \leq \end{array}$$

4 Embedding of a theory solver into an SMT solver

- **Answer isConsistent()**

Checks the so far received constraints for consistency. The answer can either be **true**, if the set of the so far received constraints is consistent, or **false**, if it is inconsistent, and **unknown**, if the theory solver cannot reason about it.

- **void pushBacktrackPoint()**

Pushes a backtrack point to the stack of backtrack points.

- **void popBacktrackPoint()**

Pops a backtrack point from the stack of backtrack points and undoes everything which has been done after adding that backtrack point.

- **vector< vector< unsigned > > getReasons()**

Returns one or more reasons for the inconsistency of the constraints. A reason is an infeasible subset of the so far received constraints. An element of a reason is a number i and means that the i th received constraint is an element of this infeasible subset.

4.2 Example: Embedding a SMT-RAT composition into OpenSMT

OpenSMT is an open source SMT solver and supports the embedding of additional theory solver. A detailed instruction of how to extend OpenSMT by a theory solver can be found on <http://verify.inf.usi.ch/opensmt>. Unfortunately, OpenSMT does not yet support quantifier free nonlinear real arithmetic (QF_NRA). For the following we assume that it supports QF_NRA and create a theory solver using the interfaces OpenSMT provides.

4.2 Example: Embedding a SMT-RAT composition into OpenSMT

```
class TSolver
{
    void    inform      ( Enode * );
    bool    assertLit   ( Enode * );
    bool    check       ( bool );

    void    pushBktPoint ( );
    void    popBktPoint  ( );
    bool    belongsToT   ( Enode * );
    void    computeModel ( );

    vector < Enode * > & explanation;
    vector < Enode * > & deductions;
    vector < Enode * > & suggestions;
}
```

OpenSMT gives the following explanation:

"**inform** is used to communicate the existence of a new T-atom to the T-solver. **assertLit** asserts a (previously informed) T-atom in the T-solver with the right polarity; it may also perform some cheap form of consistency check. **check** determines the T-satisfiability of the current set of asserted T-atoms. **pushBktPoint** and **popBktPoint** are used respectively to save and to restore the state of the T-solver, in order to cope with backtracking within the SAT-Solver. **belongsToT** is used to determine if a given T-atom belongs to the theory solved by the T-solver. Finally **computeModel** forces T-solver to save the model (if any) inside Enode's field.

Three vectors, **explanation**, **deductions**, **suggestions**, are shared among the T-solvers, and they are used to simplify the communication of conflicts, T-atoms to be deduced and "suggested" T-atoms. Suggestions are atoms consistent with the current state of the T-solver, but that they cannot be directly deduced. Suggestions are used to perform decisions in the SAT-Solver." ¹

We derive a theory solver and extend it by the three additional members

- **mpManager**, being a pointer to the theory solver object,
- **mReceivedEnodes**, needed to reconstruct which enode corresponds to which constraint,
- **mBackendsBacktrackpoints**, to store the backtrack points.

¹<http://verify.inf.usi.ch/opensmt/build-your-solver>

4 Embedding of a theory solver into an SMT solver

Figure 4.1 shows the resulting code for the header of the theory solver within `OpenSMT`. The constructor of the embedded theory solver is shown in Figure 4.2. Note, that instead of `NRATSolver` any other `SMT-RAT`-composition can be used. Figure 4.3 shows the method which informs the theory solver about the existence of a constraint. Here, we first get the string representation of the constraint, which is in prefix notation, and pass it to the `NRATSolver`. Note, that, for a reason we are not aware of, the return value has to be `lbool` and not, as mentioned on the webpage, `void`. The method `assertLit`, given by Figure 4.4, takes again the string representation of the given literal/constraint and adds it to the `NRATSolver`. The return value is `false` if the constraint is inconsistent and `true` otherwise. Note, that the polarity indicates whether the literal is positive or negative. A negativ fulfilled literal implies that the aforementioned inversion of the constraint has to hold. The implementation of the method `pushBacktrackPoint` shown by Figure 4.5 is straightforward.

Figure 4.6 shows the implementation of the method `popBacktrackPoint`. It first empties the explanation vector and then removes all stored enodes, which have been added after the backtrack point to remove. Finally, it pops the backtrack point within the theory solver used by `OpenSMT` and calls `popBacktrackPoint` of the `NRATSolver`. The check procedure given in Figure 4.7 calls `isConsistent` of the `NRATSolver`. The easiest case is if the return value is `true`. Then we know that the set of the so far added constraints is consistent and return `true`. If it determines inconsistency, that is `isConsistent` returns `false`, we extract one reason for the inconsistency and store the corresponding literals in the explanation vector. This is why we need to store the so far received literals. Finally, we return `false`. Note, that we might obtain more than one reason. However, `OpenSMT` does not support more than one reason and we discard the other reasons. The case of getting the answer `unknown` is unfortunately not supported by `OpenSMT`.

4.2 Example: Embedding a SMT-RAT composition into OpenSMT

```
#include "TSolver.h"
#include <smt-rat/smt-rat.h>

class NRASolver : public OrdinaryTSolver
{
public:
    NRASolver( const int
               , const char *
               , SMTConfig &
               , Egraph &
               , SStore &
               , vector< Enode * > &
               , vector< Enode * > &
               , vector< Enode * > & );

    ~NRASolver ( );

    lbool    inform          ( Enode * );
    bool     assertLit       ( Enode *, bool = false );
    void     pushBacktrackPoint ( );
    void     popBacktrackPoint ( );
    bool     check           ( bool );
    bool     belongsToT      ( Enode * );
    void     computeModel     ( );
#ifdef PRODUCE_PROOF
    Enode *  getInterpolants  ( );
#endif
private:
    Manager*      mpManager      ;
    vector< Enode* > mReceivedEnodes ;
    vector< unsigned > mBacktrackPoints;
};
```

Figure 4.1: The header implementation of a theory solver used by OpenSMT.

```

NRASolver::NRASolver( const int          i
                      , const char *      n
                      , SMTConfig &        c
                      , Egraph &           e
                      , SStore &           t
                      , vector< Enode * > & x
                      , vector< Enode * > & d
                      , vector< Enode * > & s )
    : OrdinaryTSolver ( i, n, c, e, t, x, d, s )
{
    mpManager      = new NRATSolver()      ; \\ This could also be
                                           \\ your theory solver
                                           \\ composition.

    mReceivedEnodes = vector< Enode* >() ;
    mBacktrackPoints = vector< unsigned >();
}

```

Figure 4.2: The constructor of a theory solver used by OpenSMT.

```

lbool NRASolver::inform(Enode *e)
{
    (void)e;
    assert(e);
    assert(belongsToT(e));

    ostringstream stream;
    stream << e;

    mpManager->inform( stream.str(), false );

    return l_Undef;
}

```

Figure 4.3: The method to inform the theory solver used by OpenSMT about a constraint.

4.2 Example: Embedding a SMT-RAT composition into OpenSMT

```
bool NRASolver::assertLit( Enode *e, bool reason )
{
    (void)e;
    (void)reason;
    assert(e);
    assert(belongsToT(e));

    mReceivedEnodes.push_back( e );

    ostringstream stream;
    stream << e;

    return mpManager->addConstraint( stream.str()
                                     , false
                                     , e->getPolarity()==l_True );
}
```

Figure 4.4: The method to assert a literal in a theory solver used by OpenSMT.

```
void NRASolver::pushBacktrackPoint()
{
    mBacktrackPoints.push_back( mReceivedEnodes.size() );
    mpManager->pushBacktrackPoint();
}
```

Figure 4.5: The method to push a backtrack point in a theory solver used by OpenSMT.

```
void NRASolver::popBacktrackPoint( )
{
    explanation.clear();

    while( mBacktrackPoints.back()<mReceivedEnodes.size() )
    {
        mReceivedEnodes.pop_back();
    }

    mBacktrackPoints.pop_back();

    mpManager->popBacktrackPoint();
}
```

Figure 4.6: The method to pop a backtrack point in a theory solver used by OpenSMT.


```
bool NRASolver::check( bool _complete )
{
    (void)_complete;

    switch( mpManager->isConsistent() )
    {
        case TS_True: return true;
        case TS_False:
        {
            vector< vector< unsigned > > reasons
                = mpManager->getReasons();
            vector< unsigned >::const_iterator pos
                = reasons.back().begin();
            while( pos!=reasons.back().end() )
            {
                explanation.push_back( mReceivedEnodes.at( *pos ) );
                ++pos;
            }
            return false;
        }
        case TS_Unknown: assert( false );
        default: assert( false );
    }
}
```

Figure 4.7: The method to check for consistency in a theory solver used by OpenSMT.

5 Implementing further modules

In this chapter we explain how to implement further modules. One principle of the design of SMT-RAT is that you can take advantage of all the features it provides, but you do not have to.

5.1 Quickstart

Figure 5.1 shows the most rudimentary way of how a module can be implemented. Furthermore, there are a few more things which must be added to SMT-RAT:

1. Add the module type `MT_MyModule` to the `enum ModuleType` of all module types in `src/ModuleType.h`.
2. The file `src/modules/Modules.h` has to include the header of the new module, i.e. add `"#include "MyModule.h"` to `modules/Modules.h`.
3. Extend the mapping of module types to module factories in the constructor of the class `Manager`:

```
addModuleType
(
    MT_MyModule,
    new StandardModuleFactory<MyModule>()
);
```

4. Add a proposition to the class `src/Formula.h` which expresses for a formula that your module cannot solve it, that is returns `unknown` or calls its backend (for further details to this see Section 5.4). Each existing proposition is represented by a bit in a `bitset` within `Formula`, currently of size 64. Hence, the proposition to add has to take an unused bit. If all bits are used the `bitset` must be extended. An example of how to add this proposition is given by Figure 5.2.
5. Extend the method `Formula::notSolvableBy` by the case that `moduleType` is `MT_MyModule` returning the before created proposition as shown in Figure 5.3
6. Extend the method `Formula::resetSolvableByModuleFlags...`
7. Finally, the file `src/CMakeLists.txt` has to be updated. Add to the set of `lib_modules_headers` `"src/modules/MyModule.h"` and to the set of `lib_modules_src` `"src/modules/MyModule.cpp"`.

```
#include "../Module.h"

class MyModule : public Module
{
public:
    MyModule( Manager* const _manager )
        : Module( _Manager )
    { mModuleType = MT_MyModule; }

    ~MyModule( ) { }
};
```

Figure 5.1: The most rudimentary way of a module can be implemented.

```
..
    static const Condition PROP_CANNOT_BE_SOLVED_BY_MYMODULE
        = Condition().set( 53, 1 );
..
```

Figure 5.2: Example of how to add a proposition to `Formula.h`.

The generated module now uses the straightforward implementation of the super-class `Module`. This includes many useful methods to read and manipulate the data in the module and the interfaces which can be extended such that they contain the actual implementation of your module.

5.2 Members of a modul

Before we explain the virtual interfaces containing the actual implementation of your modul we give a small description of the members of `Module`.

- `vector< unsigned > mBackTrackPoints`
Stores the backtrack points.
- `vector< Formula > mInfeasibleSubsets`
Stores the infeasible subsets of the so far received constraints, if the module determined their inconsistency before. An object of the type `Formula` currently consists of a conjunction of constraints and stores which of the propositions, defined in Chapter 6, hold.
- `Manager* const mpManager`
A pointer to the manager which stores this module. By this pointer the modul can obtain its backend, which gets assigned by the manager as it implements the strategy.

```

void Formula::notSolvableBy( ModuleType _moduleType )
{
    switch( _moduleType )
    {
        ..
        case MT_MyModule:
        {
            mPropositions |= PROP_CANNOT_BE_SOLVED_BY_MYMODULE;
            break;
        }
        ..
    }
}

```

Figure 5.3: How to extend the method `Formula::notSolvableBy` in `Formula.cpp`.

- `ModuleType mModuleType`
The type of this module.
- `vector< Module* > mBackend`
The pointer to the module's backend. The backend is stored in the manager.
- `vector< unsigned > mBackendBacktrackpoints`
The backtrack points in the backend.
- `bool mBackendUptodate`
A flag which indicates that the current backend considers the so far passed constraints.
- `Formula* mpReceivedConstraints`
The vector of the so far received constraints.
- `Formula* mpPassedConstraints`
The vector of the so far passed constraints.

5.3 Virtual interfaces

In the following we explain the virtual interfaces of the module, showing what their implementation in `Module` already provides and how they should be extended in your module.

5.3.1 Informing about a constraint

```
void inform( Constraint* _constraint )
```

Informs the module about the existence of the given constraint. This information might be useful for the module. The implementation in the `Module` does just nothing by calling this method.

```
void MyModule::inform( Constraint* _constraint )
{
    /*
     * Write the implementation here.
     */
}
```

Figure 5.4: Example showing how to implement the method `inform`

5.3.2 Adding a constraint

```
bool addConstraint( Constraint* _constraint )
```

Adds a constraint to the vector of the so far received constraints of the module. If the module can already decide whether the given constraint is not consistent, it returns `false` otherwise `true`. This is exactly how it is implemented in `Module`, unless it does no kind of consistency check on the given constraint, that is it always adds the given constraint to the vector of the so far received constraints and returns `true`. Note, that implementing your own module might need some initialization in this method, but you should always call `Module::addConstraint` at the beginning of your implementation of `addConstraint`.

```
bool MyModule::addConstraint( Constraint* _constraint )
{
    Module::addConstraint( _constraint );
    /*
     * Write the implementation here.
     */
}
```

Figure 5.5: Example showing how to implement the method `addConstraint`

5.3.3 Checking for consistency

```
Answer isConsistent()
```

Implements the actual consistency check of the constraints this module has so far received. There are three options how this module can answer: it can determine that these constraints are consistent and returns **true**, or it determines inconsistency and returns **false**. Otherwise, it returns **unknown**. If this module has a backend, it can call it to check the consistency of any set of constraints. How to run the backend is shown in Section 5.4. The method `Module::isConsistent` does nothing more than directly calling its backend on the so far received constraints. A module has also the opportunity to reason about the conflicts occurred, if it determines inconsistency. For this purpose it can store several infeasible subsets of the set of so far received constraints, which can be accessed as seen before.

```
Answer MyModule::isConsistent( )
{
    /*
     * Write the implementation here.
     */
}
```

Figure 5.6: Example showing how to implement the method `isConsistent`

5.3.4 Pushing a backtrack point

```
void pushBacktrackPoint()
```

Adds a backtrack point to the stack of backtrackpoints. This is already implemented by `Module::pushBacktrackPoint`. It might be necessary to store the current state of the datastructure used in your module.

```
void MyModule::pushBacktrackPoint( )
{
    Module::pushBacktrackPoint();
    /*
     * Write the implementation here.
     */
}
```

Figure 5.7: Example showing how to implement the method `pushBacktrackPoint`

5.3.5 Popping a backtrack point

```
void popBacktrackPoint()
```

Removes the last backtrack point from the stack of backtrack points and undoes everything which has been done after adding that backtrack point.

5 Implementing further modules

`Module::popBacktrackPoint` already removes the corresponding constraints from the vector of the so far received constraints, applies `popBacktrackPoint` in its backend until it does not contain any constraint created as a consequence of a removed received constraint, and removes all infeasible subsets containing removed received constraints. If you implement `popBacktrackPoint` make sure that everything stored additionally, which depends on the removed received constraints, gets deleted. Call `Module::popBacktrackPoint` afterwards.

```
void MyModule::popBacktrackPoint( )
{
    /*
     * Write the implementation here.
     */
    Module::popBacktrackPoint();
}
```

Figure 5.8: Example showing how to implement the method `popBacktrackPoint`

5.4 Running the backend module

Fortunately, there is no need to manage the addition of constraints to the backend, or pushing and popping its backtrack points. This would be even more involved as we do allow changing the backend if it is appropriate (more details to this are explained in Chapter 6). Running the backend is done in two steps:

1. Fill up the vector of the so far passed constraints.
2. Call `runBackend()`.

The first step is a bit more tricky, as we need to know which received constraints led to a passed constraint. For this purpose a modul contains a mapping from a passed constraint to one or more sets of received constraints. We give a small example by showing the behavior of one of the modules we provide, the `SimplifierModule`. Let us assume that this modul has so far received the following constraints:

$$c_0 : x \leq 0, \quad c_1 : x \geq 0, \quad c_2 : x = 0$$

`SimplifierModule` combines the first two constraints c_0 and c_1 to $c_3 : x = 0$. Then it combines c_3 with c_2 to $c_5 : x = 0$. Afterwards it calls its backend on the only remaining constraint, that is the passed constraints contain only $c_5 : x = 0$. The mapping of c_5 to the received constraints it stems from is

$$c_5 \mapsto (\{c_0, c_1\}, \{c_2\}).$$

The mapping is maintained automatically and offers two methods to add constraints to the vector of the passed constraints:

5.4 Running the backend module

- `bool addReceivedConstraintToPassedConstraints(unsigned _pos)`

Adds the constraint at the given position of the vector of the so far received constraints to the vector of the passed constraints. The mapping to its *original constraints* contains only the constraint at the given position in the vector of received constraints. It returns `true` if a constraint has been added, which only occurs if the given position is indeed within the range of the vector of received constraints and this constraint has not yet been added.

- `bool addConstraintToPassedConstraints`
(
 `Constraint* _constraint,`
 `vector< set< Constraint* > > _origins`
)

Adds the given constraint to the vector of the passed constraints. It is mapped to the given original constraints `_origins`. It returns `true` if a constraint has been added, which only occurs if the given constraint has not yet been added. Note, that `_origins` must contain only members of the vector of so far received constraints. We do not check this by reason of efficiency.

6 Composing a solver

SMT-RAT contributes a toolbox for composing an SMT compliant solver for real arithmetic, that means it is incremental, supports backtracking and provides reasons for inconsistency. The resulting solver is either a fully operative SMT solver for real arithmetic, which can be applied directly on `.smt2`-files, or a theory solver, which can be embedded into an SMT solver in order to extend its supported logics by real arithmetic.

We are talking about composition and toolbox, as SMT-RAT contains implementations of many different mathematical approaches to tackle real arithmetic, each of them embedded in a module with uniform interfaces. These modules form the tools in the toolbox and it is dedicated to a user how we use them for solving a real arithmetic formula. We provide a self-explanatory graphical user interface (GUI) for defining a graph-based *strategy* specifying which module(s) should be applied on which formula, taking into account the modules we have already involved.

In the following of this chapter we firstly give a formal definition of a strategy, followed by a brief introduction to the existing modules equipped with an estimation of their input-based performances. Afterwards we illustrate how to use the GUI for composing a strategy.

6.1 The strategy

For the comprehension of this section we presume the reader to be familiar with the concept of `runBackends(...)`, which we introduced in Chapter ??.

From the viewpoint of a user, who composes a solver with SMT-RAT, the composition starts solving a real arithmetic formula φ by invoking a consistency check of a module M on φ . Most of the modules simplify φ to an equisatisfiable real arithmetic formula φ' or need to know the consistency of a real arithmetic subproblem φ' , which it cannot solve itself. In that case, M calls `runBackends(...)`, which runs the consistency check of a backend module M' on φ' . The strategy specifies M' according to, firstly, the structure of φ' and, secondly, the call hierarchy of the consistency checks of modules leading to this request.

Therefore, the strategy is a directed tree $T := (V, E)$ with a set V of module instances as nodes and $E \subseteq V \times \Omega \times V$, where Ω is a set of conditions on real arithmetic formulas. Initially, the method `isConsistent(φ_{rcv})` of the module instance given by the root of the strategy is called, where φ_{rcv} is a conjunction of real arithmetic formulas. Whenever a module instance $m \in V$ calls `runBackends(φ_{pass})`, `isConsistent(φ_{pass})` of each module m' is called, for which an edge $(m, \omega, m') \in E$ exists such that ω holds for φ_{pass} , and passes the results back to m . A condition on real arithmetic formulas are arbitrary Boolean combinations of properties of a real arithmetic formula as follows.

6 *Composing a solver*

- Propositions about the Boolean structure of the formula. In our case it is always a conjunction of positive literals, e.g. whether it is a conjunction of real arithmetic constraints only or whether it is in conjunctive normal form.
- Propositions about the constraints in the formula, e.g. whether the formula contains (does not contain) (only) equations/(strict inequalities).
- Propositions about the polynomials compared by the constraints, e.g., whether they are linear, univariate, contain more than n variables or have a degree smaller than k .

6.2 Specifying a strategy with the GUI